

Customer Purchasing Intention

Vivek Teja

5/23/2020

R Markdown

This is an R Markdown document. Markdown is a simple formatting syntax for authoring HTML, PDF, and MS Word documents. For more details on using R Markdown see <http://rmarkdown.rstudio.com>.

When you click the **Knit** button a document will be generated that includes both content as well as the output of any embedded R code chunks within the document. You can embed an R code chunk like this:

```
summary(cars)
```

```
##      speed      dist
##  Min.   : 4.0    Min.   : 2.00
## 1st Qu.:12.0    1st Qu.: 26.00
## Median :15.0    Median : 36.00
## Mean   :15.4    Mean   : 42.98
## 3rd Qu.:19.0    3rd Qu.: 56.00
## Max.   :25.0    Max.   :120.00
```

Including Plots

You can also embed plots, for example:



Note that the `echo = FALSE` parameter was added to the code chunk to prevent printing of the R code that generated the plot.

```
# loading packages
library(dplyr)
```

```
##
## Attaching package: 'dplyr'

## The following objects are masked from 'package:stats':
##
##   filter, lag

## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union
```

```
library(ggplot2)
library(pROC)
```

```
## Type 'citation("pROC")' for a citation.
```

```
##
## Attaching package: 'pROC'
```

```
## The following objects are masked from 'package:stats':  
##  
##     cov, smooth, var
```

```
library(glmnet)
```

```
## Warning: package 'glmnet' was built under R version 3.6.3
```

```
## Loading required package: Matrix
```

```
## Loaded glmnet 3.0-2
```

```
library(readr)  
online_shoppers_intention <- read_csv("online_shoppers_intention.csv")
```

```
## Parsed with column specification:  
## cols(  
##   Administrative = col_double(),  
##   Administrative_Duration = col_double(),  
##   Informational = col_double(),  
##   Informational_Duration = col_double(),  
##   ProductRelated = col_double(),  
##   ProductRelated_Duration = col_double(),  
##   BounceRates = col_double(),  
##   ExitRates = col_double(),  
##   PageValues = col_double(),  
##   SpecialDay = col_double(),  
##   Month = col_character(),  
##   OperatingSystems = col_double(),  
##   Browser = col_double(),  
##   Region = col_double(),  
##   TrafficType = col_double(),  
##   VisitorType = col_character(),  
##   Weekend = col_logical(),  
##   Revenue = col_logical()  
## )
```

```
#                               /**** Data preperation ****/  
# reading data and creating dataframe  
  
df = online_shoppers_intention
```

```
# dataframe dimensions and features  
glimpse(df)
```

```
## Observations: 12,330  
## Variables: 18  
## $ Administrative      <dbl> 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, ...  
## $ Administrative_Duration <dbl> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...  
## $ Informational        <dbl> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...  
## $ Informational_Duration <dbl> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...
```

```
## $ ProductRelated      <dbl> 1, 2, 1, 2, 10, 19, 1, 0, 2, 3, 3, 16, 7, 6, ...
## $ ProductRelated_Duration <dbl> 0.000000, 64.000000, 0.000000, 2.666667, 627....
## $ BounceRates         <dbl> 0.200000000, 0.000000000, 0.200000000, 0.0500...
## $ ExitRates           <dbl> 0.200000000, 0.100000000, 0.200000000, 0.1400...
## $ PageValues          <dbl> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...
## $ SpecialDay          <dbl> 0.0, 0.0, 0.0, 0.0, 0.0, 0.0, 0.4, 0.0, 0.8, ...
## $ Month               <chr> "Feb", "Feb", "Feb", "Feb", "Feb", "Feb", "Feb", "Fe...
## $ OperatingSystems    <dbl> 1, 2, 4, 3, 3, 2, 2, 1, 2, 2, 1, 1, 1, 2, 3, ...
## $ Browser             <dbl> 1, 2, 1, 2, 3, 2, 4, 2, 2, 4, 1, 1, 1, 5, 2, ...
## $ Region              <dbl> 1, 1, 9, 2, 1, 1, 3, 1, 2, 1, 3, 4, 1, 1, 3, ...
## $ TrafficType         <dbl> 1, 2, 3, 4, 4, 3, 3, 5, 3, 2, 3, 3, 3, 3, 3, ...
## $ VisitorType         <chr> "Returning_Visitor", "Returning_Visitor", "Re...
## $ Weekend             <lgl> FALSE, FALSE, FALSE, FALSE, TRUE, FALSE, FALS...
## $ Revenue             <lgl> FALSE, FALSE, FALSE, FALSE, FALSE, FALSE, FAL...
```

```
# dataframe summary
summary(df)
```

```
## Administrative      Administrative_Duration Informational
## Min.   : 0.000      Min.   : 0.00      Min.   : 0.0000
## 1st Qu.: 0.000      1st Qu.: 0.00      1st Qu.: 0.0000
## Median : 1.000      Median : 7.50      Median : 0.0000
## Mean   : 2.315      Mean   : 80.82      Mean   : 0.5036
## 3rd Qu.: 4.000      3rd Qu.: 93.26      3rd Qu.: 0.0000
## Max.   :27.000      Max.   :3398.75      Max.   :24.0000
## Informational_Duration ProductRelated      ProductRelated_Duration
## Min.   : 0.00      Min.   : 0.00      Min.   : 0.0
## 1st Qu.: 0.00      1st Qu.: 7.00      1st Qu.: 184.1
## Median : 0.00      Median : 18.00      Median : 598.9
## Mean   : 34.47      Mean   : 31.73      Mean   : 1194.8
## 3rd Qu.: 0.00      3rd Qu.: 38.00      3rd Qu.: 1464.2
## Max.   :2549.38      Max.   :705.00      Max.   :63973.5
## BounceRates          ExitRates          PageValues          SpecialDay
## Min.   :0.000000      Min.   :0.00000      Min.   : 0.000      Min.   :0.00000
## 1st Qu.:0.000000      1st Qu.:0.01429      1st Qu.: 0.000      1st Qu.:0.00000
## Median :0.003112      Median :0.02516      Median : 0.000      Median :0.00000
## Mean   :0.022191      Mean   :0.04307      Mean   : 5.889      Mean   :0.06143
## 3rd Qu.:0.016813      3rd Qu.:0.05000      3rd Qu.: 0.000      3rd Qu.:0.00000
## Max.   :0.200000      Max.   :0.20000      Max.   :361.764      Max.   :1.00000
## Month              OperatingSystems      Browser              Region
## Length:12330      Min.   :1.000      Min.   : 1.000      Min.   :1.000
## Class :character      1st Qu.:2.000      1st Qu.: 2.000      1st Qu.:1.000
## Mode  :character      Median :2.000      Median : 2.000      Median :3.000
##                      Mean   :2.124      Mean   : 2.357      Mean   :3.147
##                      3rd Qu.:3.000      3rd Qu.: 2.000      3rd Qu.:4.000
##                      Max.   :8.000      Max.   :13.000      Max.   :9.000
## TrafficType        VisitorType          Weekend          Revenue
## Min.   : 1.00      Length:12330      Mode :logical      Mode :logical
## 1st Qu.: 2.00      Class :character      FALSE:9462      FALSE:10422
## Median : 2.00      Mode  :character      TRUE :2868       TRUE :1908
## Mean   : 4.07
## 3rd Qu.: 4.00
## Max.   :20.00
```

```
# missing value analysis
sapply(df, function(col) sum(is.na(col)))
```

```
##      Administrative Administrative_Duration      Informational
##      0              0              0
## Informational_Duration      ProductRelated ProductRelated_Duration
##      0              0              0
##      BounceRates      ExitRates      PageValues
##      0              0              0
##      SpecialDay      Month      OperatingSystems
##      0              0              0
##      Browser      Region      TrafficType
##      0              0              0
##      VisitorType      Weekend      Revenue
##      0              0              0
```

```
# creating list of numerical and categorical variables
features_numerical = c('Administrative', 'Administrative_Duration', 'Informational', 'Informational_Duration')
features_categorical = c('OperatingSystems', 'Browser', 'Month', 'Region', 'TrafficType', 'VisitorType', 'Weekend')
```

```
# creating dataframes with numerical and categorical features
df_numerical = df[features_numerical]
df_categorical = df[features_categorical]
head(df)
```

```
## # A tibble: 6 x 18
##   Administrative Administrative_Duration Informational Informational_Duration ProductRelated
##   <dbl>          <dbl>          <dbl>          <dbl>          <dbl>
## 1         0            0            0            0            1
## 2         0            0            0            0            2
## 3         0            0            0            0            1
## 4         0            0            0            0            2
## 5         0            0            0            0           10
## 6         0            0            0            0           19
## # ... with 13 more variables: ProductRelated_Duration <dbl>, BounceRates <dbl>,
## #   ExitRates <dbl>, PageValues <dbl>, SpecialDay <dbl>, Month <chr>,
## #   OperatingSystems <dbl>, Browser <dbl>, Region <dbl>, TrafficType <dbl>,
## #   VisitorType <chr>, Weekend <lgl>, Revenue <lgl>
```

```
# modifying the datatypes
df_categorical$Revenue = ifelse(df_categorical$Revenue == TRUE, 1, 0)
df_categorical$Weekend = ifelse(df_categorical$Weekend == TRUE, 1, 0)
df_categorical$TrafficType = as.factor(df_categorical$TrafficType)
df_categorical$Region = as.factor(df_categorical$Region)
df_categorical$OperatingSystems = as.factor(df_categorical$OperatingSystems)
df_categorical$Browser = as.factor(df_categorical$Browser)
glimpse(df_categorical)
```

```
## Observations: 12,330
## Variables: 8
## $ OperatingSystems <fct> 1, 2, 4, 3, 3, 2, 2, 1, 2, 2, 1, 1, 1, 2, 3, 1, 1, 1...
## $ Browser <fct> 1, 2, 1, 2, 3, 2, 4, 2, 2, 4, 1, 1, 1, 5, 2, 1, 1, 1...
```

```
## $ Month          <chr> "Feb", "Feb", "Feb", "Feb", "Feb", "Feb", "Feb", "Fe...
## $ Region         <fct> 1, 1, 9, 2, 1, 1, 3, 1, 2, 1, 3, 4, 1, 1, 3, 9, 4, 1...
## $ TrafficType    <fct> 1, 2, 3, 4, 4, 3, 3, 5, 3, 2, 3, 3, 3, 3, 3, 3, 4...
## $ VisitorType    <chr> "Returning_Visitor", "Returning_Visitor", "Returning...
## $ Weekend        <dbl> 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1...
## $ Revenue        <dbl> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0...
```

```
# feature engineering to reduce classes in categorical variables
```

```
df_categorical$OperatingSystems = ifelse(df_categorical$OperatingSystems %in% c('1','2','3'), df_categ
df_categorical$Browser = ifelse(df_categorical$Browser %in% c('1','2'), df_categorical$Browser, 'Other'
df_categorical$Region = ifelse(df_categorical$Region %in% c('1','3'), df_categorical$Browser, 'Other')
df_categorical$TrafficType = ifelse(df_categorical$TrafficType %in% c('12','17','18'), df_categorical$
```

```
# creating dummy variables
```

```
df_dummy = fastDummies::dummy_cols(df_categorical[c(-7,-8)], remove_first_dummy =
```


```
TRUE)
```

```
# replacing the factor variables with dummies
```

```
df_categorical_dummy = cbind(df_dummy[,c(-1:-6)], df_categorical[,c(7,8)])
```

```
# checking for multicollinearity
```

```
corrplot::corrplot(cor(df_numerical))
```



```


# feature engineering to remove multicollinearity
df_numerical$Mean_Administrative_Duration = df_numerical$Administrative_Duration / df_numerical$Adminis
df_numerical$Mean_Informational_Duration = df_numerical$Informational_Duration / df_numerical$Informati
df_numerical$Mean_ProductRelated_Duration = df_numerical$ProductRelated_Duration / df_numerical$Product
df_numerical[is.na(df_numerical)] = 0
df_numerical_fe = df_numerical[,c(-1:-7)]

```

```

# checking for multicollinearity
corrplot::corrplot(cor(df_numerical_fe))

```



```

# combining numerical and categorical dataframes
df_cleaned = cbind(df_numerical_fe, df_categorical_dummy)

```

```

#                               /**** Modelling ****/
set.seed(1)
# train test split - 80:20
train_index = sample(1:dim(df_cleaned)[1], dim(df_cleaned)[1]*0.8, replace = FALSE)
df_train = df_cleaned[train_index, ]
df_test = df_cleaned[-train_index, ]

```

```

# base logit model
model = glm(Revenue~., data=df_train, family='binomial')
summary(model)

```

```
##
```

```
## Call:
## glm(formula = Revenue ~ ., family = "binomial", data = df_train)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -6.0578  -0.4742  -0.3451  -0.1605   3.3005
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    -1.151e+01  5.354e+02  -0.021  0.982854
## ExitRates      -1.924e+01  1.811e+00 -10.624 < 2e-16 ***
## PageValues      8.105e-02  2.678e-03  30.267 < 2e-16 ***
## SpecialDay     -5.986e-02  2.571e-01  -0.233  0.815881
## Mean_Administrative_Duration  5.146e-04  6.387e-04   0.806  0.420404
## Mean_Informational_Duration  7.504e-04  3.892e-04   1.928  0.053880 .
## Mean_ProductRelated_Duration  1.050e-03  8.480e-04   1.238  0.215684
## OperatingSystems_2  1.010e-01  1.619e-01   0.624  0.532576
## OperatingSystems_3 -1.943e-01  1.783e-01  -1.090  0.275836
## OperatingSystems_Other -9.043e-02  1.862e-01  -0.486  0.627274
## Browser_2       -4.519e-02  2.058e-01  -0.220  0.826202
## Browser_Other    1.094e-01  2.086e-01   0.524  0.600096
## Month_Dec       -5.696e-01  2.048e-01  -2.781  0.005420 **
## Month_Feb      -1.583e+00  6.383e-01  -2.480  0.013154 *
## Month_Jul      -5.655e-02  2.504e-01  -0.226  0.821314
## Month_June     -1.421e-01  3.059e-01  -0.464  0.642322
## Month_Mar      -5.135e-01  2.020e-01  -2.543  0.011006 *
## Month_May      -5.276e-01  1.952e-01  -2.703  0.006875 **
## Month_Nov       6.701e-01  1.834e-01   3.655  0.000257 ***
## Month_Oct       5.923e-02  2.267e-01   0.261  0.793861
## Month_Sep     -6.841e-03  2.385e-01  -0.029  0.977114
## Region_2        9.713e-02  1.854e-01   0.524  0.600293
## Region_Other    3.433e-02  1.640e-01   0.209  0.834184
## TrafficType_2   -7.920e-01  5.595e+02  -0.001  0.998871
## TrafficType_Other  9.778e+00  5.354e+02   0.018  0.985429
## VisitorType_Other -6.386e-01  6.455e-01  -0.989  0.322476
## VisitorType_Returning_Visitor -1.710e-01  9.269e-02  -1.845  0.064999 .
## Weekend        1.260e-01  7.868e-02   1.601  0.109327
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 8606.4  on 9863  degrees of freedom
## Residual deviance: 5848.4  on 9836  degrees of freedom
## AIC: 5904.4
##
## Number of Fisher Scoring iterations: 12
```

```
y_pred_prob = predict(model, df_test, type='response')
y_pred = ifelse(y_pred_prob>0.5, 1, 0)
mean(df_test$Revenue == y_pred)
```

```
## [1] 0.8941606
```



```
ROC = roc(df_test$Revenue, y_pred)
```

```
## Setting levels: control = 0, case = 1
```

```
## Setting direction: controls < cases
```

```
auc(ROC)
```

```
## Area under the curve: 0.6808
```


```
# lasso regression
```

```
X = as.matrix(df_train[,c(-28)])
```

```
y = as.matrix(df_train['Revenue'])
```

```
cv_lasso <- cv.glmnet(X, y, family='binomial', alpha=1, standardize=TRUE, nfolds=5, type.measure='auc')
```

```
plot(cv_lasso)
```



```
model_lasso = glmnet(X, y, alpha=1, standardize=TRUE, lambda=cv_lasso$lambda.min)
coef(model_lasso)
```

```
## 28 x 1 sparse Matrix of class "dgCMatrix"
##                                     s0
## (Intercept)                      0.118966878
## ExitRates                       -0.091458482
```

```
## PageValues          0.006872808
## SpecialDay          .
## Mean_Administrative_Duration .
## Mean_Informational_Duration .
## Mean_ProductRelated_Duration .
## OperatingSystems_2 .
## OperatingSystems_3 .
## OperatingSystems_Other .
## Browser_2 .
## Browser_Other .
## Month_Dec .
## Month_Feb .
## Month_Jul .
## Month_June .
## Month_Mar .
## Month_May .
## Month_Nov          0.008137810
## Month_Oct .
## Month_Sep .
## Region_2 .
## Region_Other .
## TrafficType_2 .
## TrafficType_Other .
## VisitorType_Other .
## VisitorType_Returning_Visitor .
## Weekend .
```

```
# logit model with feature selection
```

```
model_fe = glm(Revenue~ExitRates+PageValues+Month_Nov, data=df_train, family='binomial')
summary(model_fe)
```

```
##
## Call:
## glm(formula = Revenue ~ ExitRates + PageValues + Month_Nov, family = "binomial",
##      data = df_train)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -6.0072  -0.4559  -0.3669  -0.1779   3.5229
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  -2.073834   0.066342  -31.26  <2e-16 ***
## ExitRates    -20.647091   1.766114  -11.69  <2e-16 ***
## PageValues     0.081613   0.002656   30.72  <2e-16 ***
## Month_Nov     1.061094   0.070801   14.99  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 8606.4  on 9863  degrees of freedom
## Residual deviance: 5913.9  on 9860  degrees of freedom
## AIC: 5921.9
```

```
##
## Number of Fisher Scoring iterations: 6

y_pred_prob = predict(model_fe, df_test, type='response')
y_pred = ifelse(y_pred_prob>0.5, 1, 0)
mean(df_test$Revenue == y_pred)
```

```
## [1] 0.892944
```

```
ROC = roc(df_test$Revenue, y_pred)
```

```
## Setting levels: control = 0, case = 1
```

```
## Setting direction: controls < cases
```

```
auc(ROC)
```

```
## Area under the curve: 0.6789
```

```
# finding the best threshold
threshold = seq(0.2, 0.5, by=0.05)
accuracy = c()
auc = c()
```

```
for (i in 1:7){
  y_pred = ifelse(y_pred_prob>threshold[i], 1, 0)
  accuracy[i] = mean(df_test$Revenue == y_pred)
  auc[i] = auc(roc(df_test$Revenue, y_pred))
}
```

```
## Setting levels: control = 0, case = 1
```

```
## Setting direction: controls < cases
```

```
## Setting levels: control = 0, case = 1
```

```
## Setting direction: controls < cases
```

```
## Setting levels: control = 0, case = 1
```

```
## Setting direction: controls < cases
```

```
## Setting levels: control = 0, case = 1
```

```
## Setting direction: controls < cases
```

```
## Setting levels: control = 0, case = 1
```

```
## Setting direction: controls < cases


## Setting levels: control = 0, case = 1

## Setting direction: controls < cases

## Setting levels: control = 0, case = 1

## Setting direction: controls < cases
```

```
df_threshold_results = data.frame(threshold, accuracy, auc)
plot(threshold, accuracy, type='b', main='Threshold vs Accuracy')
```



```
# optimized logit model
model_final = glm(Revenue~ExitRates+PageValues+Month_Nov, data=df_train, family='binomial')
summary(model_final)
```

```
##
## Call:
## glm(formula = Revenue ~ ExitRates + PageValues + Month_Nov, family = "binomial",
##      data = df_train)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
```

```
## -6.0072 -0.4559 -0.3669 -0.1779 3.5229
##
## Coefficients:
##             Estimate Std. Error z value Pr(>|z|)
## (Intercept) -2.073834  0.066342 -31.26  <2e-16 ***
## ExitRates   -20.647091  1.766114 -11.69  <2e-16 ***
## PageValues    0.081613  0.002656  30.72  <2e-16 ***
## Month_Nov     1.061094  0.070801  14.99  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 8606.4  on 9863  degrees of freedom
## Residual deviance: 5913.9  on 9860  degrees of freedom
## AIC: 5921.9
##
## Number of Fisher Scoring iterations: 6
```

```
y_pred_prob = predict(model_final, df_test, type='response')
y_pred = ifelse(y_pred_prob>0.275, 1, 0)
mean(df_test$Revenue == y_pred)
```

```
## [1] 0.9018654
```

```
ROC = roc(df_test$Revenue, y_pred)
```

```
## Setting levels: control = 0, case = 1
```

```
## Setting direction: controls < cases
```

```
auc(ROC)
```

```
## Area under the curve: 0.7664
```

```
# results matrix
models = c('Base logit model', 'Logit model with Lasso regression', 'Optimized logit model')
features = c(27, 3, 3)
accuracy = c(0.89, 0.89, 0.90)
auc = c(0.68, 0.68, 0.77)
df_results = data.frame(models, features, accuracy, auc)
df_results
```

```
##              models features accuracy  auc
## 1      Base logit model      27    0.89 0.68
## 2 Logit model with Lasso regression      3    0.89 0.68
## 3      Optimized logit model      3    0.90 0.77
```